

Classification I

Lecture 4

Manuel Balsera

IEOR, Columbia University

February 9, 2018

Outline

- 1 Introduction
 - Introduction
- 2 Density Estimation
 - Density Kernels
- 3 K-Neighbors
 - K-Neighbors
 - Examples
- 4 Naive Bayes Classifier
 - Naive Bayes
 - Examples
- 5 Conclusions
- 6 Bibliography

Road Map

- today we start our journey through ML classification algorithms.
- First we try a **brute force** algorithm with very weak assumptions about data: **K-Neighbors**.
 - We will show it can work very well given **enough data**.
 - **doomed by the curse of dimensionality.**
 - We will follow [1]Chapter 4 of Duda et al for K-Neighbors.
 - Chapter 13 of [2]The information Retrieval Book is a good reference on Naive Bayes.
- We will then try making **drastic** assumptions about data: **Naive Bayes**

If data representation is adequate, it may work remarkably well.
- We will use two examples as a guide:
 - 1 Image recognition: MNIST
 - 2 Text Classification: C50

Bayes Classifier

- Assume y is a categorical random variable taking $k = 1, \dots, K$ classes.
- If we know the conditional probability $p(y = k|x)$ we can implement the

Bayes Classifier

$$\hat{k} = \arg \max_k p(y = k|x) \quad (1)$$

- The error rate of the Bayes classifier is the **Bayes Error Rate**.
- The Bayes error rate is the lowest possible error rate for any classifier.
- We will in practice use $\hat{p}(y, x)$ and approximation to the true conditional probability
- In this class we study two approximations to \hat{p} and their classifiers

Density Estimation

Let's consider a random variable $S \in \mathcal{S}$ with density $p(s)$.

The probability that a sample s_i falls in region $R \subset \mathcal{S}$ is

$$\theta = \int_R ds p(s) \quad (2)$$


This is a Bernoulli random variable. Given N samples $\{s_i\}$ the maximum likelihood estimate of θ is

$$\hat{\theta} = \frac{\hat{N}_1}{N} = \frac{k}{N} \quad (3)$$


where $\hat{N}_1 = k$ is the number of samples s_i that fell inside the region R . If the region R has a small volume V centered around s_0 we have the

Monte Carlo Density Estimator


$$\hat{p}(s_0) \approx \frac{k}{N} \frac{1}{V} \quad (4)$$




(a) $N = 50$ per class, $\hat{N}_B = 11$,
 $\hat{N}_R = 5$, $\hat{p}_B = 0.88\%$, $\hat{p}_R = 0.4\%$



(c) $N = 500$ per class, $\hat{N}_B = 89$,
 $\hat{N}_R = 59$, $\hat{p}_B = 0.68\%$, $\hat{p}_R = 0.48\%$



(b) Upper Right $\hat{N}_B = 4$, $\hat{N}_R = 0$,
 $\hat{p}_B = 1.28\%$, $\hat{p}_R = 0.0\%$



(d) Upper Right $\hat{N}_B = 28$, $\hat{N}_R = 14$,
 $\hat{p}_B = 0.90\%$, $\hat{p}_R = 0.44\%$

Table 1: Count estimates for two classes

Convergence of Density Estimator

- We want a small V so that $\int_V ds p(s) \approx p(s_0)V$
- But if V is small the number of “hits” k will be small, our estimate $\hat{\theta}$ will be bad.

Consider a sequence of experiments with $N \rightarrow \infty$, where we make V_N and k_N depend on N . Define

$$\hat{p}_N(s_0) = \frac{k_N}{N} \frac{1}{V_N} \quad (5)$$

As $N \rightarrow \infty$ for convergence we need

$$\begin{aligned} V_N &\rightarrow 0 \\ k_N &\rightarrow \infty \\ \frac{k_N}{N} &\rightarrow 0 \end{aligned} \quad (6)$$

Parzen Windows and K-nearest neighbours

We can assure convergence two ways:

Parzen Window Estimation fix $V_N = \frac{1}{\sqrt{N}}$, determine k_N from data.

- we control **bias** explicitly through V_N
- When density is low, k_N is small: **high variance**
- When density is high, k_N is large: **low variance**
- For fixed N can be problematic (k could be zero).

K-neighbors Estimation fix $K_N = \sqrt{N}$, determine V_N from data.

- Control **variance** explicitly. Method is **adaptive**.
- When density is high, V_N will be small: **low bias**
- When density is low V_N is high: **high bias**

Top Hat Kernel

Assume $V_N = h_N^D$ is a D-dimensional hyper-cube centered on s_0 .
Define the **top hat function**

$$\varphi(u) = \begin{cases} 1 & |u_d| < \frac{1}{2} \quad d = 1, \dots, D \\ 0 & \text{otherwise} \end{cases} \quad (7)$$

$$\varphi\left(\frac{s_i - s_0}{h_N}\right) = \begin{cases} 1 & s_i \text{ in } s_0 \text{ hypercube} \\ 0 & s_i \text{ out of hypercube} \end{cases} \quad (8)$$

$$k_N = \sum_i^N \varphi\left(\frac{s_i - s_0}{h_N}\right) \quad (9)$$

$$\hat{p}_N(s_0) = \frac{1}{N} \sum_{i=1}^N \frac{1}{V_N} \varphi\left(\frac{s_i - s_0}{h_N}\right) \quad (10)$$

Density Kernels

We can generalize the definition of φ to any probability density

Kernel Density

Given a kernel density

$$\begin{aligned}\varphi(u) &\geq 0 \\ \int du \varphi(u) &= 1\end{aligned}\tag{11}$$

we have

$$\hat{p}_h(s_0) = \frac{1}{N} \sum_i^N \frac{1}{h^D} \varphi\left(\frac{s_i - s_0}{h}\right)\tag{12}$$

Density estimate is a **convolution** of a kernel density φ over data samples.

Common Kernels

Top Hat $\varphi(u) = \frac{1}{2}$ if $|u| < 1$.

Gaussian $\varphi(u) \propto e^{-\frac{u^2}{2}}$.

Epanechnikov $\varphi(u) \propto 1 - u^2$ for $|u| < 1$.

Exponential $\varphi(u) \propto e^{-|u|}$.

Linear $\varphi(u) \propto 1 - |u|$ for $|u| < 1$.

Cosine $\varphi(u) \propto \cos \frac{\pi u}{2}$ for $|u| \leq 1$.




Figure 1: Adapted from: Jake Vanderplas - sklearn

Bandwidth and Bias-Variance Tradeoff

- h controls the bias/variance tradeoff. Is called *bandwidth* or *kernel width*.
- Kernels are similarity measures φ decreases with u .
- Expected value over all samples of $p(s_0)$ is a **convolution**

$$\mathbb{E}(\hat{p}(s_0)) = \int_S ds p(s) \varphi\left(\frac{s - s_0}{h}\right) \quad (13)$$

- Larger h implies lower variance (see Section 4.3 of [Duda's book](#))

$$\text{Var}((\hat{p}(s_0))) \leq \frac{\sup(\varphi) \hat{p}(s_0)}{Nh^D} \quad (14)$$

Course of Dimensionality

- Assume $\mathcal{S} = [0, 1]^D$ is the unit hyper-cube. $p(s) \approx 1$ uniform.
- With resolution $h \approx 0.1$, $V = 10^{-D}$. We need $N \approx 10^D$ samples to estimate \hat{k} for all s .
- Become impractical to sample \mathcal{S} with any precision.
- Sometimes $P(S)$ is concentrated in a manifold (hyper-surface) of **effective dimension** $D_{\text{eff}} \ll D$ and sampling is still possible.




Figure 2: A $D_{\text{eff}} = 2$ manifold embedded in \mathbb{R}^3 . Source: Wikipedia

K-Neighbors Class Probabilities

Given a categorical variable Y taking values $y = 1, \dots, L$, and N samples $s_i = \{y_i, x_i\}$ we can estimate $p(y|x)$ using bayes theorem

$$p(y|x) = \frac{p(x, y)}{p(x)} = \frac{p(y, x)}{\sum_{y'} p(y', x)} \quad (15)$$

As before, we have the Maximum Likelihood estimate

$$\hat{p}(y, x) \approx \frac{k_y}{N} \frac{1}{V} \quad (16)$$

where V is a small volume centered around x , and k_y is the number of observations i of category y that fell in volume V .

Given $K = \sum_y k_k$, the number of samples that fell on volume V we have

K Neighbors Probabilities

$$\hat{p}(y|x) \approx \frac{k_y}{K} \quad (17)$$

K-Neighbors Classifier

- If we keep the volume V fixed, and use a kernel φ we have **kernel classification**.
- If we fix K and let V be determined by the data $\{x_i\}$ we have a **nearest neighbors** classifier.
- To minimize classification error we select the class with higher probability.

K-Neighbors Classifier

Given a point x , predict the most frequent class y along the K closest neighbors x_i on the training data.

Properties of K-Neighbors Classifier

- Training is trivial: store all samples $\{y_i, x_i\}$
- Prediction is expensive $O(DN)$: Search through all samples. There are [speed up techniques](#).
- The parameter K controls the Bias-Variance trade-off.
- A special case is the **Nearest Neighbor** classifier.
- As $N \rightarrow \infty$ Nearest Neighbor classifier error is bounded by *twice* the Bayes error rate.

Example Density Estimation

Let's sample 100 points from a **Mixed Gaussian** distribution

- $S \sim \mathcal{N}(0, 1)$ with probability 0.3
- $S \sim \mathcal{N}(5, 1)$ with probability 0.7

We set the bandwidth to $h = \frac{1}{2}$






Figure 3: Adapted from: Jake Vanderplas—sklearn


Ellipse Classification Problem




(a) Nearest Neighbors classifier with $N = 100$. Accuracy 88%.



(b) 31st nearest neighbors classifier with $N = 100$. Accuracy 60%.



(c) Nearest Neighbors classifier with $N = 1,000$. Accuracy 90%.



(d) 31st nearest neighbors classifier with $N = 1,000$. Accuracy 92.5%.

MNIST Classification


- Images are represented as $R \times C = 28 \times 28$ matrices $F_{r,c}$.
- Feature vector X_d is the linearized image pixels $d = 1, \dots, R \times C$

$$X_{i,C*r+w} = F_{i,r,c} \quad (18)$$

- Distance is regular Euclidean L_2 distance for X_d
- Nearest neighbors has **96.9% accuracy** in test sample.
- Prediction is **slow**, can only classify 8 images per second.
- Prediction requires the whole training sample: consumes disk and memory.

MNIST Sample Preparation

- This two images share no pixels (very far in L_2 distance).



- MNIST images have been carefully prepared (centered).
- If we train with digits shifted ± 3 pixels, out of sample performance **drops to 47%**
- Images **must be carefully prepared** for K -neighbors to do well.

C50 Text Classification

- We classify the **C50** Reuters text documents using our pre-build feature vectors.
- We use **cosine** distance (L_2 distance would be dominated by document length).

| Features | $K = 1$ | $K = 5$ |
|----------|---------|---------|
| Set | 59.4% | 61.4% |
| Count | 53.9% | 52.7% |
| Tf-Idf | 54.3% | 54.6% |

Figure 4: K-Neighbors out of sample classification accuracy on the C50 Documents

- There are 50 balanced classes. Random guessing has 2% accuracy.
- **Set** binary features outperform: Once a word is included in a text there seem to be little extra information if it re-appears.

Application: Protein 3D Structure Prediction




Figure 5: DNA binding protein Hist1 sequence alignment

- A protein is composed of sequence of amino acids.
- There are 22 choices, represented by letters A, R, N, etc...
- We define a distance $d(A, B)$ between pairs of amino acids




Figure 6: Hist1 3D structure.

- We wish to predict a protein 3D structure from its sequence.
- After aligning two sequences we can define a distance based on the sum of paired amino acid distances
- We can then use **nearest neighbors** to predict 3D structure.

Naive Bayes Assumption

Curse of Dimensionality

The complexity of describing an **arbitrary function** $f(x)$ for $x \in \mathcal{X} \subset \mathbb{R}^D$ **grows exponentially** in D .

- in a classification problem we need to estimate $P(Y = y|x)$ for $y = 1, \dots, K$, a set of K separate D -dimensional functions.
- We need some **simplifying assumption**.
- Simplest approach is to assume different dimensions in x are **conditionally independent** given the class label y .

Naive Bayes Assumption

$$P(x|y) = \prod_d^D P(x_d|y) \quad (19)$$

Naive Bayes Classifier

We really need $P(y|x)$. Using Bayes theorem

$$P(y|x) = \frac{P(x, y)}{P(x)} = \prod_d^D P(x_d|y) \frac{P(y)}{P(x)} \quad (20)$$

- Very **strong** assumption. Hardly ever justified.
- Probability estimate almost always **wrong**.
- But we **hope** relative ordering of class probabilities preserved.
- Taking logs and dropping $P(x)$ (does not depend on y):

Naive Bayes Classifier

$$\hat{y}_{\text{NB}} = \arg \max_y \sum_d \log P(x_d|y) + \log P(y) \quad (21)$$

Properties of Naive Bayes

- **Drastic simplification:** Learn D 1-dimensional functions, rather than one D -dimensional function.
- Ignores all **dependencies** between features x_d .
- We can make different assumptions on $P(x_d|y)$ based on the data. For example:

Bernoulli if x_d is a binary variable

Multinomial if x_d is categorical

Gaussian if x_d is continuous and approximately Gaussian.

- It is really a **framework**: Can mix and match assumptions for different features x_d .
- **Not invariant** to equivalent re-statements of features x_d .
If x_A and x_B are independent:

Gaussian $x_A + x_B$ and $x_A - x_B$ not independent.

Bernoulli $x_A \text{ XOR } x_B$ and x_A not independent.

Gaussian Naive Bayes

Given sample data $\{y_i, x_{i,d}\}$ where

- $i = 1, \dots, N$ runs over the data samples.
- y_i is a categorical variable taking the values $k = 1, \dots, K$
- the input data is D -dimensional $d = 1, \dots, D$
- $x_{i,d}$ given y is approximately normally distributed

The Gaussian Naive Bayes assumption is

$$p(x_{i,d}|y = k) = \mathcal{N}(x_{i,d}; \hat{\mu}_{d,k}, \sigma_{d,k}^2) \quad (22)$$

where different dimensions d are independent.

Using the one-hot $z_{i,k}$ representation of the labels y_i the max likelihood estimate of $p(x_{i,d}|y_i = k)$ has parameters

$$\begin{aligned} \hat{\mu}_{d,k} &= \frac{1}{\hat{N}_k} \sum_i x_{i,d} z_{i,k} = \frac{1}{\sum_i z_{i,k}} \sum_i x_{i,d} z_{i,k} \\ \hat{\sigma}_d^2 &= \frac{1}{\hat{N}_k} \sum_i (x_{i,d} - \mu_{d,k})^2 z_{i,k} \end{aligned} \quad (23)$$

Gaussian Naive Bayes Loss Function

The marginal probability of class k is

$$\hat{\pi}_k = \frac{\hat{N}_k}{N} \quad (25)$$

The loss function for a sample $x = (x_1, \dots, x_D)$ is

$$L_k(x) = -\log p(y = k|x) = \frac{1}{2} \sum_d \left\{ \left(\frac{x_d - \hat{\mu}_{d,k}}{\hat{\sigma}_{d,k}} \right)^2 + \log (2\pi \hat{\sigma}_{d,k}^2) \right\} - \log \hat{\pi}_k \quad (26)$$

And finally, the predicted class will be

$$p(y = k) = \hat{k} = \arg \min_k L_k(x) \quad (27)$$

Geometry of Gaussian Naive Bayes Decision Boundary

At the **decision boundary** where the probability of classes k_1 and k_2 is the same must have that

$$L_{k_1}(x) = L_{k_2}(x) \quad (28)$$

or, using the explicit expression of the loss function

$$0 = \frac{1}{2} \sum_d \left\{ \left(\frac{x_d - \hat{\mu}_{d,k_1}}{\hat{\sigma}_{d,k_1}} \right)^2 - \left(\frac{x_d - \hat{\mu}_{d,k_2}}{\hat{\sigma}_{d,k_2}} \right)^2 + \log \frac{\hat{\sigma}_{d,k_1}^2}{\hat{\sigma}_{d,k_2}^2} \right\} - \log \frac{\hat{\pi}_{k_1}}{\hat{\pi}_{k_2}} \quad (29)$$

- This is a D-dimensional quadratic surface. If $D = 2$ this is a conic section: ellipse, hyperbola or parabola.
- Because there are no cross terms $x_d x_{d'}$, the conic axis must be align with the coordinate axis

Gaussian Naive Bayes Example

The following two examples have $N = 40$ data points with $D = 2$.




Figure 7: x_1 and x_2 are independent.
Accuracy $\approx 78\%$




Figure 8: x_1 and x_2 are dependent.
Accuracy $\approx 68\%$

Single Variance Approximation

Gaussian Naive Bayes approximation estimates $2D \times K + K - 1$ parameters

If data is scarce we can make the approximation that variance does not depend on class

$$\hat{\sigma}_d^2 = \frac{1}{N} \sum_k \sum_i (x_{i,d} - \hat{\mu}_{d,k})^2 z_{i,k} \quad (30)$$

The equation satisfied by the boundary between two classes is then

$$0 = \frac{1}{2} \sum_d \left\{ \left(\frac{x_d - \hat{\mu}_{d,k_1}}{\hat{\sigma}_d} \right)^2 - \left(\frac{x_d - \hat{\mu}_{d,k_2}}{\hat{\sigma}_d} \right)^2 \right\} - \log \frac{\hat{\pi}_{k_1}}{\hat{\pi}_{k_2}} \quad (31)$$

which simplifies to

$$0 = \sum_d \frac{(\hat{\mu}_{d,k_1} - \hat{\mu}_{d,k_2})}{\hat{\sigma}_d^2} \left\{ \frac{(\hat{\mu}_{d,k_1} + \hat{\mu}_{d,k_2})}{2} - x_d \right\} - \log \frac{\hat{\pi}_{k_1}}{\hat{\pi}_{k_2}} \quad (32)$$

a **linear equation** for the boundary.

Naive Bayes and Text

Text representations are very high dimensional \rightarrow Naive Bayes.

- A document is a sequence of words $D = \{w_1, \dots, w_t, \dots, w_T\}$.
- We wish to classify documents in $y = 1, \dots, K$ classes.
- **Naive Bayes Assumption:** given y , words w_t are chosen **independently** from a multinomial distribution.
- w_t is one of $1, \dots, d, \dots, V$, where V is the size of the **vocabulary**.
- probability

$$P(D|y = k) = \prod_t^T P(w_t|y = k) \quad (33)$$

We need to estimate $P(w_t|y)$ from an N document **corpus** $D_i = \{w_{i,t}\}$.

Estimating Word Probabilities

- Given count $x_{i,d}$ of word d in document i we need to estimate the probability that word d is generated in class k .
- word d is a categorical variable with D possible values.
- The number of categories is large: $D = V$ the size of vocabulary.
- As V is large most of the counts will be small: we must use

Bayesian smoothing

- The estimate of generating word d in class k is

$$\hat{P}(w = d | y = k) = \frac{\hat{n}_{k,d} + \alpha}{\hat{n}_k + \alpha V} \quad (34)$$

- The counts of word d in class k are

$$\hat{n}_{k,d} = \sum_i x_{i,d} z_{i,k} \quad (35)$$

- The total number of words in class k

$$\hat{n}_k = \sum_d \hat{n}_{k,d} \quad (36)$$

Naive Bayes Text Classifier

Given the count x_d of word $d = 1, \dots, V$ in document $D = \{w_1, \dots, w_T\}$. The probability that document D is generated in class k is

$$P(D|y = k) = \prod_t^T P(w_t|y = k)P(y = k) = \prod_d^D \hat{P}(w = d|y = k)^{x_d} \hat{P}(y = k) \quad (37)$$

Taking logs and defining

$$w_{k,d} = \log \hat{P}(w = d|y = k) = \log \frac{\hat{n}_{k,d} + \alpha}{\hat{n}_k + \alpha V}$$

$$b_k = \log \hat{P}(y = k) = \frac{\hat{N}_k}{N} \quad (38)$$

where \hat{N}_k is the number of documents in class k . The log probability is

$$\hat{L}_k = \left(\sum_d x_d w_{k,d} + b_k \right) \quad (39)$$

Naive Bayes Text Classifier

$$y_{\text{NB}} = \arg \max_k \hat{L}_k = \arg \max_k \left(\sum_d x_d w_{k,d} + b_k \right) \quad (40)$$

Naive-Bayes: MNIST

- Features are the linearized pixels $X_{i,Cr+c}$.
- $P(x_d|y)$ is probability that a pixel be on in class y .
- Gaussian model does poorly ($\approx 56\%$ accuracy): pixels have a bi-modal distribution, either on or off.
- Bernoulli approximation does much better $\approx 84\%$ (using $\alpha = 1$).
- Loadings $\log P(x_d|y)$ are easy to interpret

NB Loadings Digit 0



NB Loadings Digit 5



NB Loadings Digit 1



NB Loadings Digit 6



NB Loadings Digit 2



NB Loadings Digit 7



NB Loadings Digit 3



NB Loadings Digit 8



NB Loadings Digit 4



NB Loadings Digit 9




Naive-Bayes: Translation Invariance

- Allowing **off-center digits** degrades performance.

| Train | Test | Accuracy |
|------------|------------|----------|
| Centered | Centered | 84% |
| Centered | Off-Center | 38% |
| Off-Center | Off-Center | 55% |

- Loadings $\log P(x_d|y)$ are blurred.



- With enough blurring pixel densities will be uniform: Naive Bayes is useless (K-Neighbors still works given enough data).
- Problem is with our **features**: images are defined by **pixel correlations** not positions.

Naive-Bayes: C50 Documents

- Set Features again do better

| Features | Accuracy |
|----------|----------|
| Set | 65.8% |
| Count | 65.0% |
| Tf-Idf | 63.5% |

- Set features again do better, but by only a slight margin.
- Naive-Bayes outperform K-neighbors: Text is very high-dimensional.
- For text classification Naive Bayes is the **baseline** model.

Conclusions

- We have discussed the **simplest** Machine learning algorithms.
- They make very different assumptions about feature distribution.
 - [K-neighbors](#) : non-parametric. \mathcal{X} is continuous and has a metric.
 - [Naive Bayes](#) : conditional independence of features.
- They are both useful baselines to compare more sophisticated models to.
- They both can do remarkably well, but data must be prepared carefully.
- We have begun to explore how **features** affect performance.
- K-neighbor can be quite costly at prediction time, not always practical.
- We must find a disciplined way to choose **hyper-parameters**: K , α , text feature (*next lecture*).

Bibliography



Richard O. Duda, Peter E. Hart, and David G. Stork.

Pattern Classification (2nd Ed).

Wiley, 2001.

[http:](http://cns-classes.bu.edu/cn550/Readings/duda-etal-00.pdf)

[//cns-classes.bu.edu/cn550/Readings/duda-etal-00.pdf](http://cns-classes.bu.edu/cn550/Readings/duda-etal-00.pdf).



Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze.

Introduction to Information Retrieval.

Cambridge University Press, 2008.

[https://nlp.stanford.edu/IR-book/
information-retrieval-book.html](https://nlp.stanford.edu/IR-book/information-retrieval-book.html).